

OpenACC: Directive-based Programming For Hardware Accelerators

Programming Hardware Accelerators

Scientific Applications

Software Libraries

Drop-in Replacement
of
Relevant Functions

- Easy of use
- Performance guarantee
- Fixed interface

Compiler directives

Easy Acceleration
of
Custom Code

- Portable
- Full assistance of the compiler

Programming Languages

- Flexibility of implementation
- Steep learning curve
- Maximum performance
- Error Prone

- Not for beginners

OpenACC in a Nutshell

- Directives in
 - pragmas (C, C++) or
 - comments Fortran
- Multiple parallelism levels
- Portable across
 - Accelerators
 - NVIDIA GPUs
 - Sunway
 - PEZY-SC
 - Including CPU target
 - Compilers
 - GNU GCC 9 (OpenACC 2.6)
 - PGI, ...
- Performance
 - Often close to “native” efficiency

```
int main(void) {  
    // host code  
    #pragma acc kernels  
    {  
        // accelerated code  
    }  
    return 0;  
}
```

OpenACC Main Features

- Incremental
 - Sequential code preserved
 - Annotations limited to important parts of the code
 - Profile-based performance tuning
 - Correctness verified as needed without massive code changes
- Single Source
 - Same code for sequential and accelerated code
 - Recompile on a new platform for performance
 - Sequential code unaffected and developed in unison
- Ease of use
 - Gentle learning curve
 - Works for HPC languages: C, C++, Fortran
 - No low-level hardware experience or details required

Jacobi Iteration: C Code

```
while ( err > tol && iter < iter_max ) {  
    err=0.0;  
    for( int j = 1; j < n-1; j++) {  
        for(int i = 1; i < m-1; i++) {  
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                                A[j-1][i] + A[j+1][i]);  
  
            err = max(err, abs(Anew[j][i] - A[j][i]));  
        }  
    }  
    for( int j = 1; j < n-1; j++) {  
        for( int i = 1; i < m-1; i++ ) {  
            A[j][i] = Anew[j][i];  
        }  
    }  
    iter++;  
}
```

Iterate until convergence

Iterate across matrix elems.

Calculate new value
from neighbors

Compute max-error
for convergence

Swap input/output arrays

OpenACC Syntax Overview

```
#pragma acc <directive> <clause_1>    !$acc <directive> <clause_1>
{                                         ! block of code to accelerate
    // block of code to accelerate      !$acc end <directive>
}
```

- A **pragma** in C and C++
 - Instructs the compiler how to compile the code
 - Compilers may ignore pragma they don't understand
- Specially formatted comment is required in Fortran
 - Comments can be ignored by the compiler
- **acc** starts OpenACC directives and clauses
- **Directives** are commands for OpenACC compiler to alter the following code
- **Clauses** are augment directives with extra details and functionality

OpenACC Directives

Manage
Data
Movement

```
#pragma acc data copyin(x,y) copyout(z)
{
```

Initiate
Parallel
Execution

```
/* ... */
```

```
#pragma acc parallel
```

```
{
```

Optimize
Loop
Mappings

```
#pragma acc loop gang vector
```

```
for (i = 0; i < n; ++i) {
```

```
    z[i] = x[i] + y[i];
```

```
/* ... */
```

```
}
```

```
}
```

```
/* ... */
```

```
}
```

- Incremental
- Single source
- Interoperable
- Performance portable
- CPU, GPU, Xeon Phi

OpenACC
Directives for Accelerators

OpenACC Example: Stencil Code with DATA and KERNELS

```
#pragma acc data copy(b[0:n][0:m]) create(a[0:n][0:m])
{
  for (iter = 1; iter <= p; ++iter){
    #pragma acc kernels
    {
      for (i = 1; i < n-1; ++i){
        for (j = 1; j < m-1; ++j){
          a[i][j]=w0*b[i][j]+
            w1*(b[i-1][j]+b[i+1][j]+
              b[i][j-1]+b[i][j+1])+
            w2*(b[i-1][j-1]+b[i-1][j+1]+
              b[i+1][j-1]+b[i+1][j+1]);
        }
      }
      for( i = 1; i < n-1; ++i )
        for( j = 1; j < m-1; ++j )
          b[i][j] = a[i][j];
    }
  }
}
```



Example: Jacobi Iteration for Poisson Equation

- Poisson equation has a simple form in 2D

- $u_{xx} + u_{yy} = f(x,y)$

- Applications include

- Electricity
 - Magnetism
 - Gravity
 - Heat distribution
 - Fluid flow
 - Torsion



$$\nabla^2 f(x, y) = 0$$

- When $f(x,y)=0$ we call it Laplace equation

$$u_{xx} = \frac{\partial^2 u}{\partial x \partial x} \approx \frac{u(x+h, y) - 2u(x, y) + u(x-h, y)}{h^2}$$

Jacobi Iteration: C Code

```
while ( err > tol && iter < iter_max ) {  
    err=0.0;  
    for( int j = 1; j < n-1; j++) {  
        for(int i = 1; i < m-1; i++) {  
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                                A[j-1][i] + A[j+1][i]);  
  
            err = max(err, abs(Anew[j][i] - A[j][i]));  
        }  
    }  
    for( int j = 1; j < n-1; j++) {  
        for( int i = 1; i < m-1; i++ ) {  
            A[j][i] = Anew[j][i];  
        }  
    }  
    iter++;  
}
```

Iterate until convergence

Iterate across matrix elems.

Calculate new value
from neighbors

Compute max-error
for convergence

Swap input/output arrays

Looking for Parallelism

```
while ( err > tol && iter < iter_max ) {  
    err=0.0;  
    for( int j = 1; j < n-1; j++) {  
        for(int i = 1; i < m-1; i++) {  
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                                A[j-1][i] + A[j+1][i]);  
  
            err = max(err, abs(Anew[j][i] - A[j][i]));  
        }  
    }  
    for( int j = 1; j < n-1; j++) {  
        for( int i = 1; i < m-1; i++ ) {  
            A[j][i] = Anew[j][i];  
        }  
    }  
    iter++;  
}
```

Data dependence
between iterations

Independent loop iterations

Max reduction required

Independent loop iterations


OpenACC Parallel Directive

Generates parallelism

```
#pragma acc parallel  
{
```

When encountering the **parallel** directive, the compiler will generate 1 or more parallel **gangs**, which execute redundantly.


```
}
```



OpenACC Loop Directive

Identifies loops to run in parallel

```
#pragma acc parallel  
{  
    #pragma acc loop  
    for (i=0; i<N; ++i)  
    {  
    }  
}
```




The **loop** directive informs the compiler which loops to parallelize.

OpenACC Parallel Loop Directive

Generates parallelism and identifies loop in one directive

```
#pragma acc parallel loop  
for (i=0; i<N; ++i)  
{  
}  
}
```

The **parallel** and **loop** directives
are frequently combined into one.



Looking for Parallelism

```
while ( err > tol && iter < iter_max ) {  
    err=0.0;  
    #pragma acc parallel loop reduction(max:err)  
    for( int j = 1; j < n-1; j++) {  
        for(int i = 1; i < m-1; i++) {  
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                                A[j-1][i] + A[j+1][i]);  
            err = max(err, abs(Anew[j][i] - A[j][i]));  
        }  
    }  
  
    #pragma acc parallel loop  
    for( int j = 1; j < n-1; j++) {  
        for( int i = 1; i < m-1; i++ ) {  
            A[j][i] = Anew[j][i];  
        }  
    }  
    iter++;  
}
```

Parallelize loop on
accelerator

Parallelize loop on
accelerator

A reduction means that all $n*m$ values
for err will be reduced to just one, the max.


OpenACC Loop Directive: PRIVATE & REDUCTION

- The **private** and **reduction** clauses are not optimization clauses, they may be required for correctness.
- **Private** = a copy of the variable is made for each loop iteration
- **Reduction** = a reduction is performed on the listed variables
 - Supports: +, *, max, min, and logical operators

Building OpenACC Code with PGI Compiler

```
$ pgcc -fast -acc -ta=tesla -Minfo=all laplace2d.c
main:
40, Loop not fused: function call before adjacent loop
    Generated vector sse code for the loop
51, Loop not vectorized/parallelized: potential early exits
55, Accelerator kernel generated
    55, Max reduction generated for error
    56, #pragma acc loop gang /* blockIdx.x */
    58, #pragma acc loop vector(256) /* threadIdx.x */
55, Generating copyout(Anew[1:4094][1:4094])
    Generating copyin(A[:][:])
    Generating Tesla code
58, Loop is parallelizable
66, Accelerator kernel generated
    67, #pragma acc loop gang /* blockIdx.x */
    69, #pragma acc loop vector(256) /* threadIdx.x */
66, Generating copyin(Anew[1:4094][1:4094])
    Generating copyout(A[1:4094][1:4094])
    Generating Tesla code
69, Loop is parallelizable
```

Performance Results: Speedup (Higher is Better)



- Question:
 - Why did OpenACC had a slow down on GPUs
- Answer:
 - Very low compute/memcpy ratio

■ Intel Xeon E5-2698v3 2.3GHz (Haswell)
■ NVIDIA Tesla K40/P100

```
while ( err > tol && iter <
iter_max ) {
    err=0.0;
```

A, Anew resident on host

copy

These copies
happen every
iteration of the
outer while loop!

A, Anew resident on host

copy

```
}
```

A, Anew resident on accelerator

```
#pragma acc parallel loop
for( int j = 1; j < n-1; j++) {
    for(int i = 1; i < m-1; i++)
    {
        Anew[j][i] = 0.25 * (
            A[j][i+1] + A[j][i-1] +
            A[j-1][i] + A[j+1][i]);
        err = max(err,
abs(Anew[j][i] - A[j][i]));
    }
}
```

A, Anew resident on accelerator

Looking for Parallelism

```
while ( err > tol && iter < iter_max ) {  
    err=0.0;  
    #pragma acc parallel loop reduction(max:err)  
    for( int j = 1; j < n-1; j++) {  
        for(int i = 1; i < m-1; i++) {  
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +  
                                A[j-1][i] + A[j+1][i]);  
            err = max(err, abs(Anew[j][i] - A[j][i]));  
        }  
    }  
    #pragma acc parallel loop  
    for( int j = 1; j < n-1; j++) {  
        for( int i = 1; i < m-1; i++ ) {  
            A[j][i] = Anew[j][i];  
        }  
    }  
    iter++;  
}
```

Does CPU need data
between these loop nests?

Does CPU need data
between iterations of the
convergence loop?

Data Regions

The **data** directive defines a region of code in which GPU arrays remain on the GPU and are shared among all kernels in that region.

```
#pragma acc data
{
#pragma acc parallel loop
/* . . . */

#pragma acc parallel loop
/* . . . */
}
```

Data region

Arrays used within the data region will remain on the GPU until the end of the data region.

Data Clauses

- `copy (list)`
 - Allocates memory on GPU and copies data from host to GPU when entering region and copies data to the host when exiting region.
- `copyin (list)`
 - Allocates memory on GPU and copies data from host to GPU when entering region.
- `copyout (list)`
 - Allocates memory on GPU and copies data to the host when exiting region.
- `create (list)`
 - Allocates memory on GPU but does not copy.
- `present (list)`
 - Data is already present on GPU from another containing data region.
- `deviceptr(list)`
 - The variable is a device pointer (e.g. CUDA) and can be used directly on the device.

Array Shaping

- Compiler sometimes cannot determine size of arrays
 - Must specify explicitly using data clauses and array “shape”
- C/C++
 - `#pragma acc data copyin(a[0:nelem]) copyout(b[s/4:3*s/4])`
- Fortran
 - `!$acc data copyin(a(1:end)) copyout(b(s/4:3*s/4))`
- Note: data clauses can be used on data, parallel, or kernels

Data Regions Have Real Consequences

```
int main(int argc, char** argv) {  
    float A[1000];  
  
    #pragma acc kernels  
    for( int iter = 1; iter < 1000;  
        ++iter){  
        A[iter] = 1.0;  
    }  
  
    A[10] = 2.0;  
  
    printf("A[10] = %f", A[10]);  
}
```

A[] Copied To GPU

A[] Copied To Host

Runs on Host

Output: A[10] = 2.0

```
int main(int argc, char** argv) {  
    float A[1000];  
    #pragma acc data copy(A)  
    {  
        #pragma acc kernels  
        for( int iter = 1; iter < 1000;  
            ++iter){  
            A[iter] = 1.0;  
        } // "kernels" end here  
        A[10] = 2.0;  
    }  
  
    printf("A[10] = %f", A[10]);  
}
```

A[] Copied To GPU

Still Runs on Host

A[] Copied To Host

Output: A[10] = 1.0

Looking for Parallelism

```
#pragma acc data copy(A) create(Anew)
while ( err > tol && iter < iter_max ) {
    err=0.0;
#pragma acc parallel loop reduction(max:err)
    for( int j = 1; j < n-1; j++) {
        for(int i = 1; i < m-1; i++) {
            Anew[j][i] = 0.25 * (A[j][i+1] + A[j][i-1] +
                                A[j-1][i] + A[j+1][i]);
            err = max(err, abs(Anew[j][i] - A[j][i]));
        }
    }
#pragma acc parallel loop
    for( int j = 1; j < n-1; j++) {
        for( int i = 1; i < m-1; i++ ) {
            A[j][i] = Anew[j][i];
        }
    }
    iter++;
}
```

Copy A to/from the accelerator only when needed.

Create Anew as a device temporary.

Building the New OpenACC Code with PGI Compiler

```
$ pgcc -fast -acc -ta=tesla -Minfo=all laplace2d.c
main:
40, Loop not fused: function call before adjacent loop
    Generated vector sse code for the loop
51, Generating copy(A[:][:])
    Generating create(Anew[:][:])
    Loop not vectorized/parallelized: potential early exits
56, Accelerator kernel generated
    56, Max reduction generated for error
    57, #pragma acc loop gang /* blockIdx.x */
    59, #pragma acc loop vector(256) /* threadIdx.x */
56, Generating Tesla code
59, Loop is parallelizable
67, Accelerator kernel generated
    68, #pragma acc loop gang /* blockIdx.x */
    70, #pragma acc loop vector(256) /* threadIdx.x */
67, Generating Tesla code
70, Loop is parallelizable
```

Performance Results: Speedup (Higher is Better)

